

Prompting token overhead across paradigms (DeepSeek-V4-Flash, $K=5$)

